

Embedded systems engineering

Distributed real-time systems

David Kendall

Time is Central to Real-time and Embedded Systems

Several timing analysis problems:

- Worst-case execution time (WCET) estimation
- Estimating distribution of execution times
- Threshold property: can you produce a test case that causes a program to violate its deadline?
- Software-in-the-loop simulation: predict execution time of particular program path

ALL involve predicting an execution time property!

WCET and BCET

- Remember the simple formula for response time analysis (no blocking or jitter).

$$R_i = C_i + \sum_{j \in hp(i)} \left\lceil \frac{R_j}{T_j} \right\rceil C_j$$

- Each C_i represents the *worst-case computation time* of its task.
- Worst-case computation time** (WCET) : the longest time taken by a program code to complete its execution (assuming no blocking, jitter or interference)
- Best-case computation time** (BCET) : the shortest time taken by a program code to complete its execution (assuming no blocking, jitter or interference)
- How to obtain values for C_i ?

Calculating execution times

- Measurement

- ▶ Need to exercise great care in obtaining measurements
- ▶ Need to take care in interpreting results
- ▶ How to know if you've measured the worst (best) case?

- Analysis

- ▶ Intended to guarantee that the worst (best) case execution time is reported
- ▶ Difficult to take account of all architectural effects: pipelines, caches, speculative execution etc.

- Let's consider a diagram that illustrates the relationship between these approaches

Timing analysis of systems




Fig. 1. Basic notions concerning timing analysis of systems. The lower curve represents a subset of measured executions. Its minimum and maximum are the *minimal* and *maximal observed execution times*, respectively. The darker curve, an envelope of the former, represents the times of all executions. Its minimum and maximum are the *best-* and *worst-case execution times*, respectively, abbreviated BCET and WCET.


Figure from R.Wilhelm et al., *ACM Transactions on Embedded Computing Systems*, Vol. 7:3, pp 36:1 – 36:53, April 2008.

Measuring execution time

Accuracy, precision, and resolution

- Real-world effects introduce *uncertainty* into measurements.
- Three important features that characterise the quality of measurements are:
 - ▶ **Accuracy**: an indication of the closeness of measurements of a quantity to that quantity's actual value according to some well-defined standard.
 - ▶ **Precision**: the degree to which repeated measurements under unchanged conditions give the same results.
 - ▶ **Resolution**: the smallest change in the underlying physical quantity that produces a noticeable response in the measurement.

Distinguishing accuracy and precision



Accurate but not precise



Precise but not accurate

Figures from http://en.wikipedia.org/wiki/Accuracy_and_precision

Measurement errors

- Difficult to separate individual contributions made to error by accuracy, precision, and resolution
 - ▶ Quantifying accuracy is hard, e.g. quantifying accuracy of interval timer requires calibration of clock source with standard measurement of time
 - ▶ Use variance of measurements to quantify precision
 - ▶ Resolution usually easy to quantify, e.g. resolution of interval timer can introduce an error of ± 1 clock period.
- Other sources of error include:
 - ▶ **Perturbing** the quantity to be measured by the act of measuring it, e.g. program statements added to a program to access a timer change the behaviour of the program being measured
 - ▶ Other **non-deterministic events** may also perturb the quantity to be measured
- Useful to classify errors as either **systematic** or **random**

Methods for measuring execution time

| <i>Method</i> | <i>Typical Resolution</i> | <i>Typical Accuracy</i> | <i>Granularity</i> | <i>Difficulty of Use</i> |
|-----------------------|---------------------------|-------------------------|--------------------|--------------------------|
| stop-watch | 0.01 sec | 0.5 sec | program | easy |
| date | 0.02 sec | 0.2 sec | program | easy |
| time | 0.02 sec | 0.2 sec | program | easy |
| prof and gprof | 10 msec | 20 msec | subroutines | moderate |
| clock() | 15-30 msec | 15-30 msec | statement | moderate |
| software analyzers | 10 μ sec | 20 μ sec | subroutine | moderate |
| timer/counter chips | 0.5-4 μ sec | 1-8 μ sec | statement | very hard |
| logic or bus analyzer | 50 nsec | half μ sec | statement | hard |

Figure from David B. Stewart, Measuring execution time and real-time performance (part 1), Dr. Dobbs Journal, November 2006
([available here](#))

Defining an interval timer for FRDM-K64F

PIT Initialisation

```
void counterInit(void) {  
    /* Open the clock gate to the PIT */  
    SIM->SCGC6 |= (1u << 23);  
  
    /* Enable the clock for the PIT timers.  
       Continue to run in debug mode  
    */  
    PIT->MCR = 0u;  
  
    /* Disable the timer */  
    PIT->CHANNEL[2].TCTRL &= ~PIT_TCTRL_TEN_MASK;  
  
    /* Period p = maximum available , bus clock f = 60 MHz, v = pf - 1 */  
    PIT->CHANNEL[2].LDVAL = 0xFFFFFFFF;  
  
    /* Clear the timer interrupt flag */  
    PIT->CHANNEL[2].TFLG |= PIT_TFLG_TIF_MASK;  
  
    /* Enable interrupt on timeout */  
    PIT->CHANNEL[2].TCTRL |= PIT_TCTRL_TIE_MASK;  
  
    /* Enable the interrupt in the NVIC */  
    NVIC_EnableIRQ(PIT2_IRQn);  
}
```

Defining an interval timer ...

```
void counterStart(void) {
    /* Start the timer running */
    PIT->CHANNEL[2].TCTRL |= PIT_TCTRL_TEN_MASK;
}

uint32_t counterStop(void) {
    uint32_t counter = 0;

    /* get the time elapsed since started */
    counter = 0xFFFFFFFF - PIT->CHANNEL[2].CVAL;

    /* Disable the timer */
    PIT->CHANNEL[2].TCTRL &= ~PIT_TCTRL_TEN_MASK;
    return counter;
}

void PIT2_IRQHandler(void) {
    /* Disable the timer */
    PIT->CHANNEL[2].TCTRL &= ~PIT_TCTRL_TEN_MASK;

    /* Disable interrupt on timeout */
    PIT->CHANNEL[2].TCTRL &= ~PIT_TCTRL_TIE_MASK;

    /* Clear the timer interrupt flag to allow further timer interrupts */
    PIT->CHANNEL[2].TFLG |= PIT_TFLG_TIF_MASK;

    /* We should never get here - timer overflow */
    assert(false);
}
```

Using an interval timer

```
int main () {
    uint32_t timeElapsed = 0;

    /* Initialise the counter. Only need to do this once */
    counterInit();

    /* Reset time elapsed */
    timeElapsed = 0;

    /* Start the counter */
    counterStart();

    /* Start of code to be measured */
    lcd.locate(2,8);
    lcd.printf("Hello world!\n");
    /* End of code to be measured */

    /* Stop the counter and retrieve time elapsed */
    timeElapsed = counterStop();

    /* Display the result */
    lcd.cls();
    lcd.locate(86, 8);
    lcd.printf("%u", timeElapsed);
}
```

Questions about the use of the interval timer

- How accurate is the timer?
 - ▶ Don't know. Don't have the equipment to calibrate it exactly. Read the data sheet for the clock crystal?
- How precise is the timer?
 - ▶ Don't know. Haven't done enough experiments yet. Need to perform multiple experiments. Compute mean and variance.
- What is the resolution of the timer?
 - ▶ Speed of peripheral clock for PIT1 is configured to be 0.5 of the SystemCoreClock, which is set at 120 MHz.
- What is the longest interval of time that can be measured with this configuration before the timer counter overflows (32 bit counter)?

Measuring execution time: summary

- There are numerous techniques for measuring software execution time.
- The use of an interval timer has been described here.
- Every approach to measurement has limitations and may introduce errors.
 - ▶ Whenever you present data based on measurement, it's important to discuss the likely accuracy, precision and resolution of the measurements.
 - ▶ Take care not to introduce systematic errors into your experiments, e.g. by perturbing the software that you are trying to measure.
 - ▶ Account for random errors by repetition of experiments to establish a level of confidence in the data.
- Worst-case execution time may not be revealed by measurement – results likely to be optimistic, ie to suggest a worst-case execution time that is less than the actual worst-case execution time.

Acknowledgements

- Raimund Kirner, Y.T. Li, S. Malik, Edward Lee, Sanjit Seshia, R. Wilhelm, *The Determination of Worst-Case Execution Times: Overview of Methods and Survey of Tools*, ACM Transactions on Embedded Computing Systems (TECS) 7:3, 2008